Neural CSSR Research Notebook

Discovering Epsilon Machines with Neural Networks

Neural CSSR Lab

October 6, 2025

Abstract

This notebook explores the challenges of classical CSSR (Causal State Splitting Reconstruction) and introduces our Neural CSSR solution. We'll examine why traditional CSSR is brittle to parameter choices and how neural networks can provide a more robust approach to discovering epsilon machines from sequential data.

1 What We're Building: Epsilon Machine Discovery

1.1 The Epsilon Machine Framework

What is an Epsilon Machine?

An **epsilon machine** (ε-machine) is the minimal predictive model of a stochastic process. It captures the **causal structure** - exactly what information from the past is needed to optimally predict the future.
Key properties:

- **Minimal**: Uses the fewest states possible
- **Predictive**: Each state has a unique probability distribution over future symbols
- **Causal**: States represent equivalence classes of histories with identical predictive power

Imagine you observe a sequence of 0s and 1s: `01101001110...`

The fundamental question: **Is there a hidden state machine generating this sequence?** If so, what does it look like?

1.2 Our Case Studies: Two Challenging Processes

In this notebook, we'll analyze two particularly challenging processes that highlight the limitations of classical CSSR:

Process 1: Even Process (Infinite Memory)

The **Even Process** tracks the parity (even/odd count) of 1s seen so far:

- **States**: Even (E) and Odd (O) - representing current parity
- **Memory requirement**: Theoretically infinite - need to count all 1s from the beginning
- **Transitions**: E $\xrightarrow{0}$ E, E $\xrightarrow{1}$ O, O $\xrightarrow{1}$ E (0 from O is forbidden)
- **Challenge**: Long-range dependencies make classical CSSR fail

Example sequence: `01110110...`

- After `0`: Even state (0 ones seen)
- After `01`: Odd state (1 one seen)
- After `011`: Even state (2 ones seen)
- After `0111`: Odd state (3 ones seen)

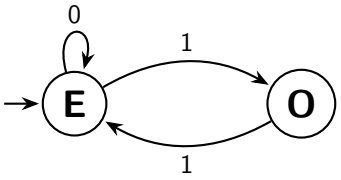

Process 2: Seven-State Human (Complex Benchmark)

The **Seven-State Human Machine** comes from human sequence prediction studies and serves as a standard benchmark:

- **States**: 7 distinct causal states with context-dependent naming (BB, AAA, AAAB, BA, BAB, BAAB, BAA)
- **Probabilities**: Precise fractions from experimental data (e.g., 15/16, 3/16, 13/16)
- **Structure**: Complex transition matrix with realistic statistical signatures
- **Challenge**: Multiple states with subtle differences require precise parameter tuning

This machine represents the complexity found in real-world sequential data and is commonly used to benchmark epsilon machine discovery algorithms.

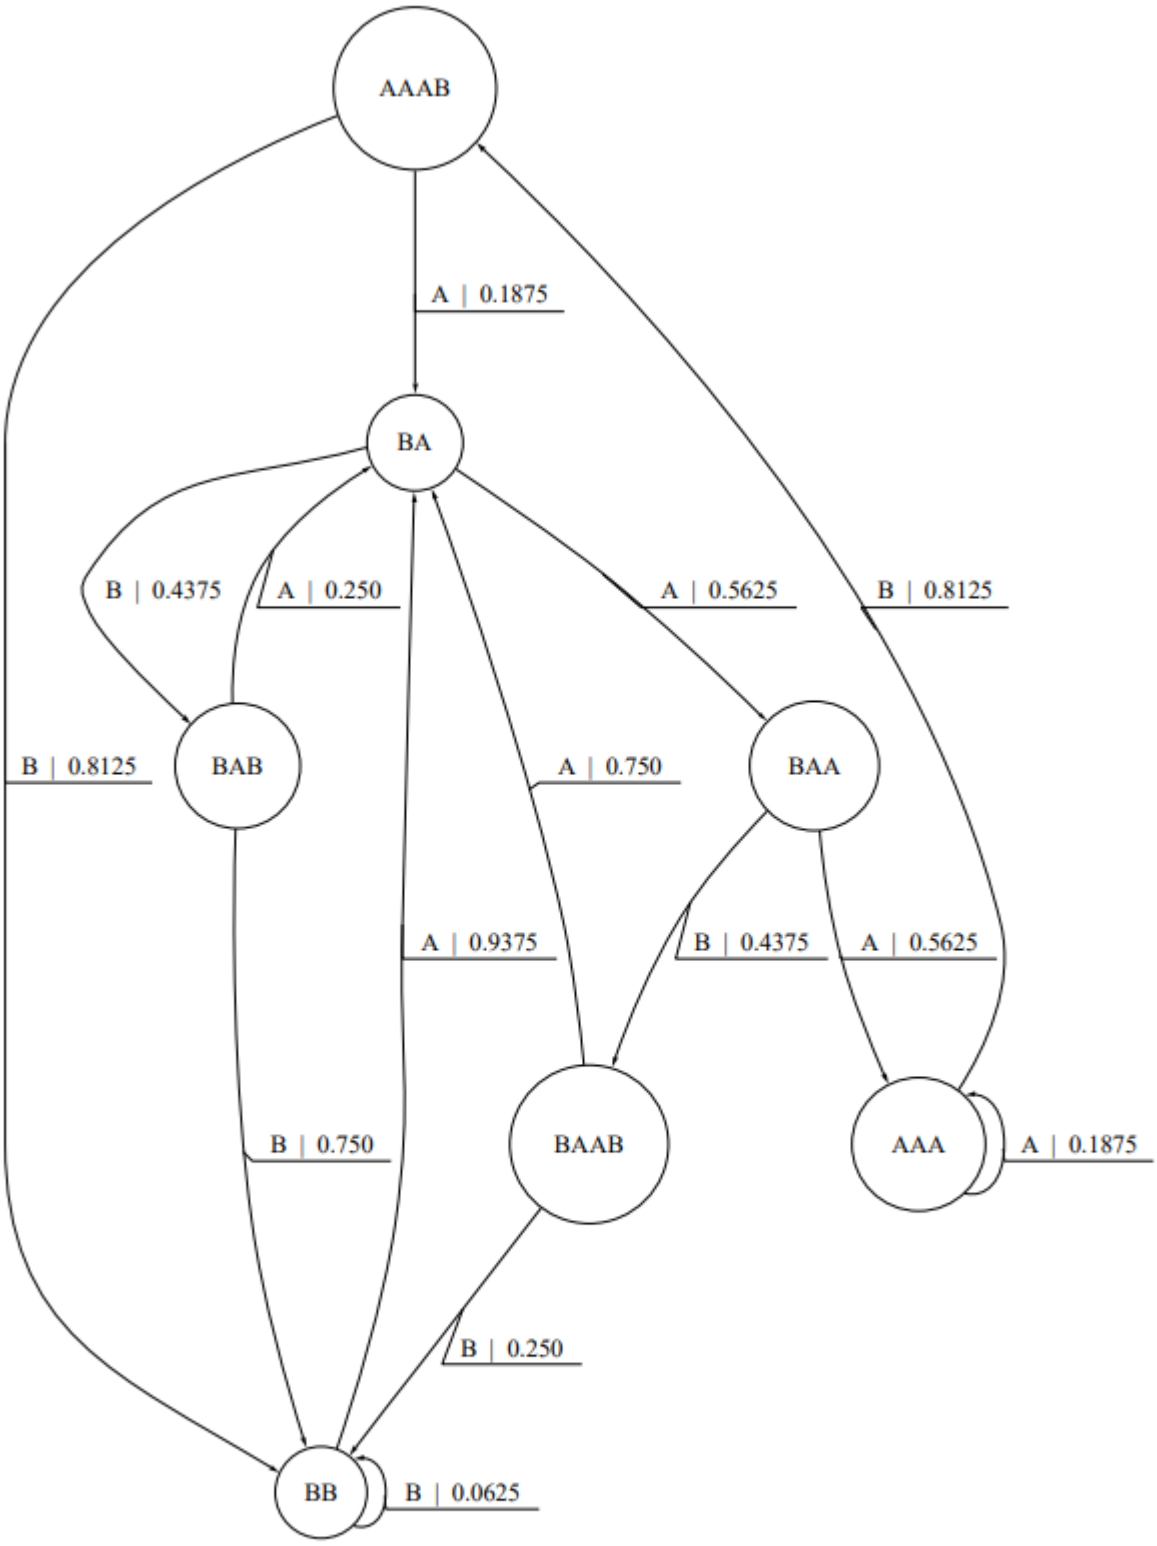

Figure 3: A seven-state process, used in [16] to study human sequence prediction. Here each state is defined by a single suffix, as indicated. All transition probabilities are multiples of 1/16.

The state machine structure is clearly shown in the diagram above, with emission probabilities detailed in the UNIFILAR properties below.

UNIFILAR Machine Properties

This 7-state machine has the **UNIFILAR** property: each (state, input symbol) leads to exactly one next state. The randomness is in the **emission probabilities**:
Key patterns from the specification:

- **BB state**: Very low B-emission probability ($\frac{1}{16} = 0.0625$), acts as absorbing hub
- **AAA/AAAB states**: High B-emission probabilities ($\frac{13}{16}$, $\frac{9}{16}$)
- **A-labeled transitions**: Create left-to-right flow through the upper chain of states
- **B-labeled transitions**: Route probability mass back to BB and close the lower loops

2 Classical CSSR: The Foundation

Classical CSSR is the gold standard algorithm for epsilon machine reconstruction. Developed in the 1990s, it's a systematic method for discovering the causal structure of stochastic processes from observational data.

2.1 The CSSR Philosophy

Core Idea of CSSR

CSSR asks: **"Which histories have the same predictive power?"**
If two different histories h_1 and h_2 have statistically indistinguishable next-symbol distributions, then they represent the same causal state:

$$P(\text{next symbol}|h_1) \approx P(\text{next symbol}|h_2) \Rightarrow h_1, h_2 \in \text{same state}$$

The algorithm incrementally builds longer histories and tests for statistical equivalence using hypothesis testing.

Example: Golden Mean Process

Consider the Golden Mean process (forbids consecutive 1s). It has a simple 2-state structure:

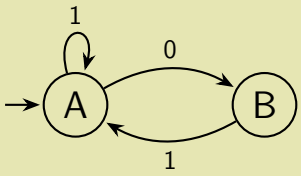

State A: Can emit 0 or 1 (but 1→1 creates the forbidden "11" pattern)
State B: Can only emit 1 (since 0 from B would allow the forbidden pattern)

However, classical CSSR suffers from significant challenges that limit its practical applicability...

2.2 The Classical CSSR Algorithm

Here's how traditional CSSR works:

Classical CSSR Algorithm

Algorithm CSSR(A, $\bar{x}$, L_{max}, α)

I. Initialization: $L \leftarrow 0$, $\Sigma \leftarrow \{\{\emptyset\}\}$

II. Sufficiency:
while $L < L_{max}$ **do**
  for each $s \in \Sigma$ **do**
    estimate $P(\widehat{X_t}|S = s)$
    for each $x \in s$ **do**
      for each $a \in A$ **do**
        estimate $p \leftarrow P(\widehat{X_t}|X_{t-L}^{t-1} = ax)$
        Test(Σ, p, ax, s, α)
      end for
    end for
  end for
  $L \leftarrow L + 1$
end while

III. Recursion:
Remove transient states from Σ
recursive $\leftarrow$ False
repeat
  recursive $\leftarrow$ True
  for each $s \in \Sigma$ **do**
    for each $b \in A$ **do**
      $x_0 \leftarrow$ first $x \in s$
      $T(s, b) \leftarrow \hat{\varepsilon}(x_0 b)$
      for each $x \in s$, $x \neq x_0$ **do**
        if $\hat{\varepsilon}(xb) \neq T(s, b)$ **then**
          create new state $s' \in \Sigma$
          $T(s', b) \leftarrow \hat{\varepsilon}(xb)$
          for each $y \in s$ such that $\hat{\varepsilon}(yb) = \hat{\varepsilon}(xb)$ **do**
            Move(y, s, s')
          end for
          recursive $\leftarrow$ False
        end if
      end for
    end for
  end for
until recursive

Test(Σ, p, ax, s, α):
if null hypothesis passes test of size α **then**
  $s \leftarrow ax \cup s$

else if restricted alternative hypothesis passes test of size α for $s^* \in \Sigma$, $s^* \neq s$ **then**
  Move(ax, s, s^*)
else
  create new state $s' \in \Sigma$
  Move(ax, s, s')
end if

Move(x, s_1, s_2):
$s_1 \leftarrow s_1 \setminus x$
re-estimate $P(\widehat{X_t}|S = s_1)$
$s_2 \leftarrow s_2 \cup x$
re-estimate $P(\widehat{X_t}|S = s_2)$

2.3 The Parameter Brittleness Problem

Classical CSSR suffers from extreme sensitivity to its two critical parameters:

Parameter Brittleness Issues

- L_{max} **(Maximum History Length)**:
  - Too small → Under-splitting (states lumped together)
  - Too large → Over-splitting (spurious states)
  - **No principled way to choose optimal value**
- α **(Significance Level)**:
  - Too high → Fails to detect real differences (under-splitting)
  - Too low → Spurious differences (over-splitting)
  - **Requires domain expertise and extensive tuning**

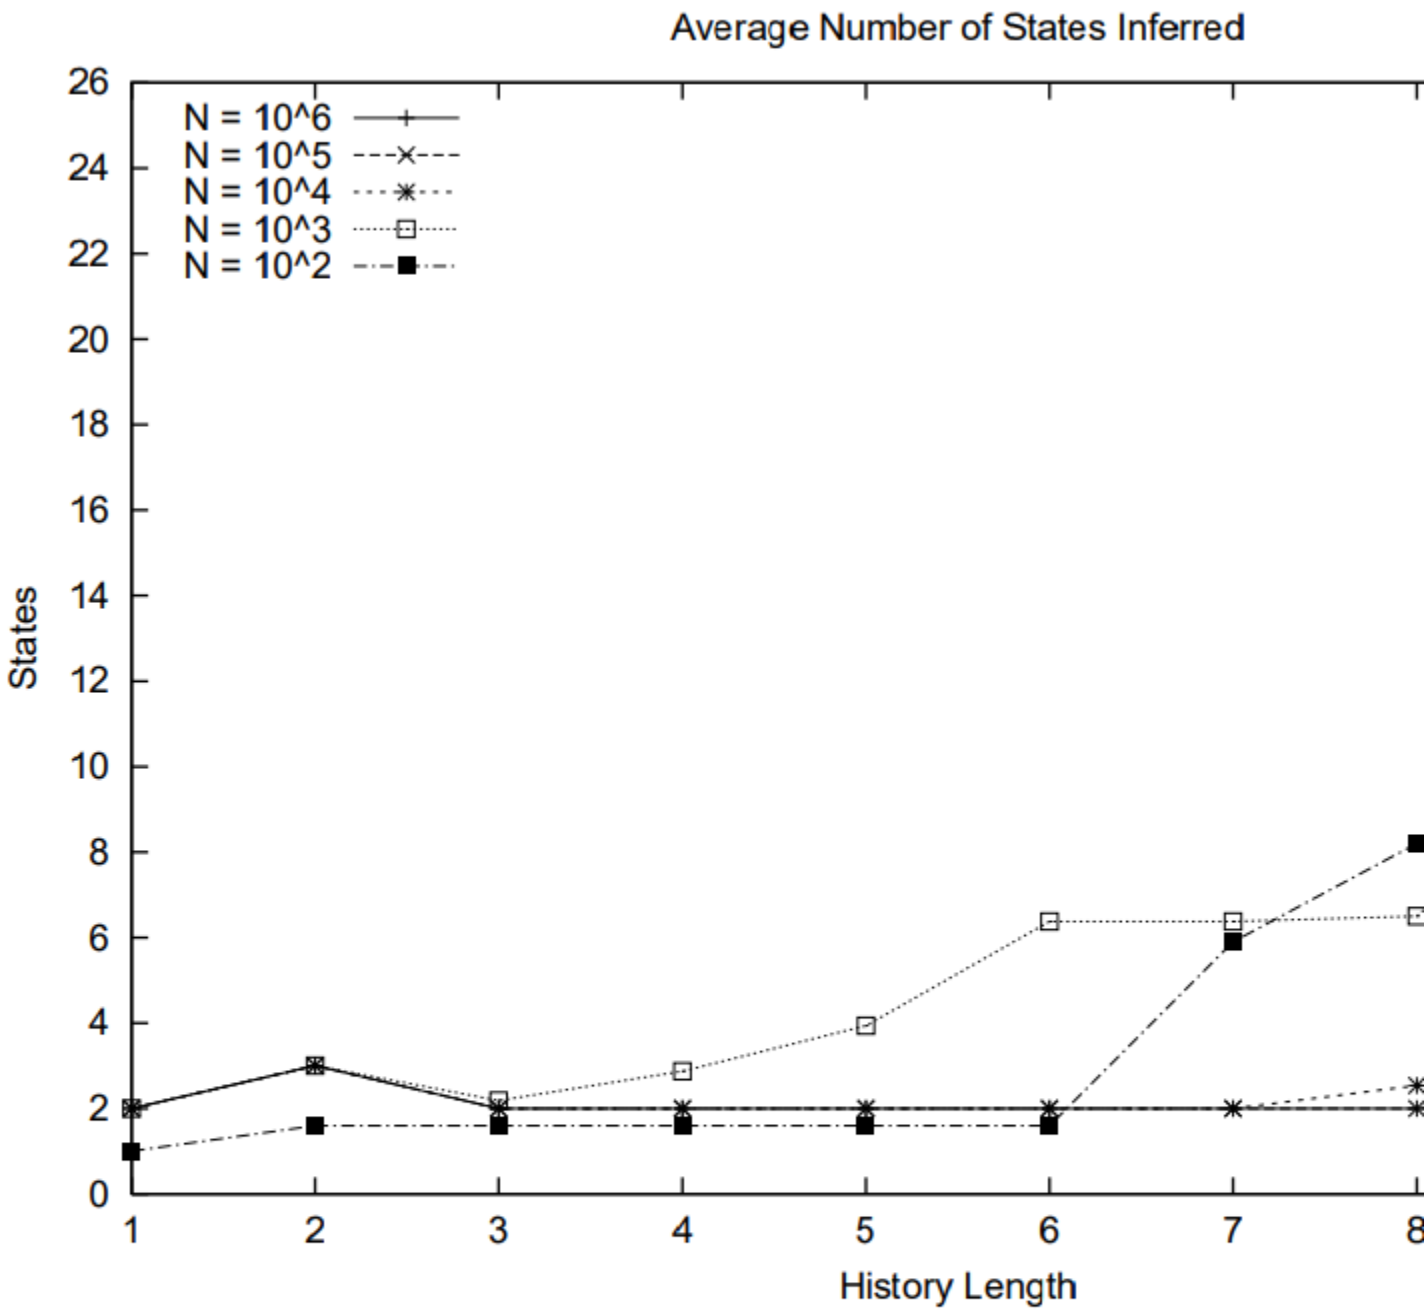

Figure 4: Number of states inferred versus
for the even process. The true number of c
2.

2.4 Why These Parameters Are So Problematic

The Goldilocks Problem

Classical CSSR requires parameters to be "just right" - but there's no systematic way to find the right values:

- **Process-dependent**: Optimal L_{max} and α vary dramatically across different processes
- **Sample-size dependent**: Parameters that work for 10K samples may fail for 100K samples
- **No cross-validation**: Can't easily validate parameter choices without ground truth

3 Neural CSSR: A Promising Solution

3.1 The Core Innovation

Neural Probability Estimation

Instead of using count-based frequency estimation, Neural CSSR replaces the probability computation step with neural network predictions:

$$P(\text{next symbol}|\text{history}) \leftarrow \text{count-based} \rightarrow \text{neural-based}$$

Key insight: Transformers can generalize across similar histories, providing reliable probability estimates even for rare or unseen sequences.

The fundamental architecture preserves CSSR's clustering mechanism while replacing its brittle counting step:

Neural CSSR Algorithm Overview

Phase I: Train transformer on sequence data
Phase II: For each history length L:
for each state s and history $h \in s$ **do**
  Compute $P(\cdot|h)$ using transformer (not counting)
  Apply same statistical tests as classical CSSR
  if distributions differ significantly **then**
    Split histories into different states
  end if
end for
Phase III: Ensure deterministic transitions (unchanged)

3.2 Implementation Results and Challenges

Experimental Findings

We implemented Neural CSSR using the `transcssr_neural_runner.py` system with encouraging preliminary results:
Promising Outcomes:

- **Reduced state count**: Neural CSSR consistently discovered fewer states than classical CSSR as L_{max} increased

- **More stable reconstruction**: Less sensitive to exact parameter choices than classical approach
- **Better generalization**: Neural probabilities handled rare histories more gracefully

Critical Challenge Identified:

- **Pseudocount scaling problem**: Converting neural probabilities to counts for χ^2 tests requires careful calibration
- **Parameter**: `pseudo_count_scale` becomes the new "brittleness point" - needs process-specific tuning
- **Statistical test mismatch**: χ^2 tests designed for integer counts, not calibrated neural probabilities

3.3 The Pseudocount Scaling Challenge

From Probabilities to Counts

Neural networks output probabilities $p \in [0, 1]$, but classical statistical tests expect integer counts. The conversion requires:

$$\text{neural probability} \xrightarrow{\text{scale}} \text{effective counts}$$

For history h with neural probability p_1 for symbol '1':

- Count for '1': $n_1 = p_1 \times \text{scale}$
- Count for '0': $n_0 = (1 - p_1) \times \text{scale}$

The problem: Choice of scale factor dramatically affects χ^2 test sensitivity!

This scaling challenge suggests that Neural CSSR's full potential may require moving beyond classical statistical tests to methods designed for continuous probability distributions.

3.4 Model Calibration

Probabilities are better represented with calibration to ensure predicted probabilities match observed frequencies. We used Platt scaling (fitting a sigmoid to validation set predictions) for our experiments. Our models were well-calibrated out of the box, but calibration doesn't hurt so we applied Platt fitting as a precautionary measure.

3.5 Future Directions

- **Distribution-aware tests**: Develop statistical tests that work directly with probability distributions rather than requiring count conversion
- **Adaptive scaling**: Automatically tune pseudocount scales based on neural probability confidence
- **Jensen-Shannon divergence**: Explore JS divergence as an alternative to χ^2 for probability-based state splitting

4 What Doesn't Work: The Hidden Representation Challenge

While our Neural CSSR approach successfully recovers epsilon machines through JS divergence analysis of predictive distributions, direct inspection of the model's hidden representations reveals fundamental limitations that highlight why we focus on output-based state discovery rather than internal representation analysis.

4.1 Highly Entangled and Fractured Representation Space

Analysis of the transformer's internal hidden states for the Golden Mean process reveals that the learned representations are highly entangled and fractured, making direct state extraction from hidden activations impractical.

Hidden State Clustering Analysis

The representation geometry analysis shows poor state separation in the hidden space:

- **Low silhouette scores**: 0.158 across all layers indicates weak clustering
- **Low participation ratio**: 1.0 suggests representations use minimal dimensionality
- **Poor state purity**: Only 66.8% purity with 27.1% normalized mutual information

Even with sophisticated clustering algorithms, the hidden representations fail to cleanly separate according to the true causal states of the process.

Redundancy Analysis: Golden Mean Process

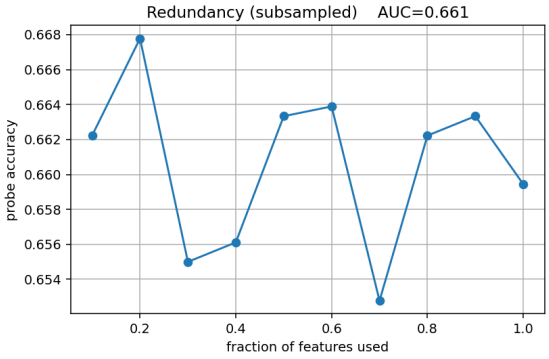

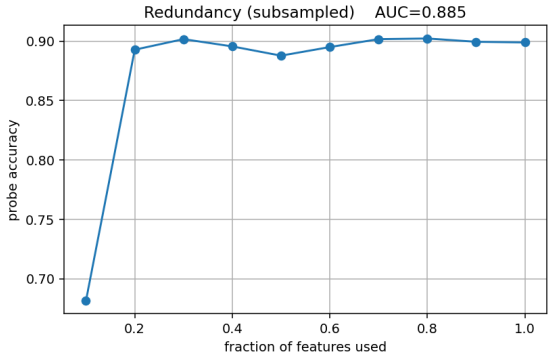

Left (Standard Golden Mean): AUC=0.661 shows minimal improvement over chance for state prediction from hidden representations. **Right (Golden Mean + Auxiliary Head)**: AUC=0.885 demonstrates that auxiliary state supervision significantly improves representation quality, but compromises the unsupervised approach.

Multi-State Process Layer Analysis

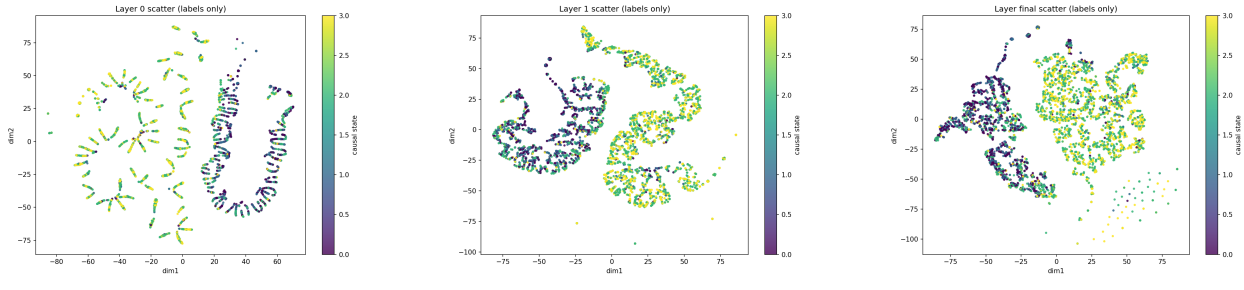

Layer 0 — Layer 1 — Final Layer: All three images show different layers from a multi-state process. Earlier layers showed more informative spaces regarding structural representation versus functional representation, but remained highly tangled and fragmented across all layers.

Multi-State Process Centroids

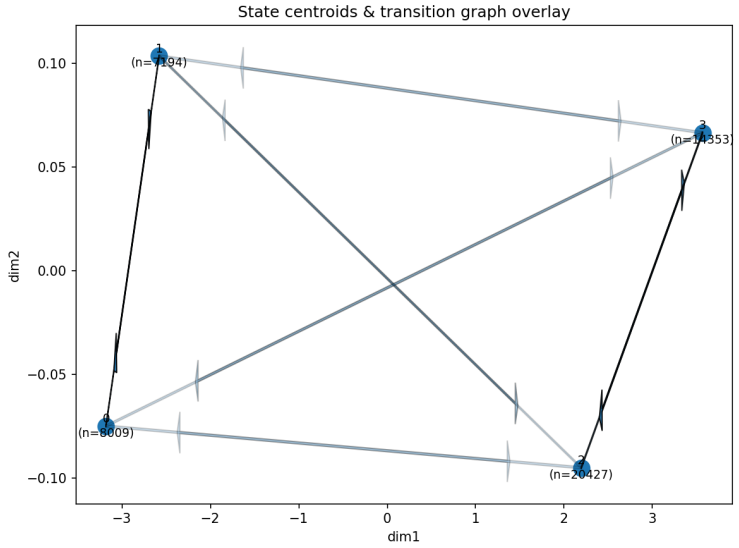

Interpretation: Centroid analysis from the multi-state process confirms the entangled nature of hidden representations, with no clear geometric structure corresponding to causal states even when examining cluster centers.

4.2 The Auxiliary State Prediction Experiment

To test whether adding explicit state supervision could improve representation quality, we conducted an experiment training an auxiliary state prediction head alongside the standard language modeling objective. While this showed some promising initial results, it was ultimately abandoned to preserve the purely unsupervised nature of our approach.

Why We Avoided Supervised State Prediction

Adding a state prediction head would compromise the core philosophy of unsupervised epsilon machine discovery:

- Requires ground-truth state labels during training

- Defeats the purpose of discovering unknown causal structure
- Would not generalize to real-world processes where true states are unknown
- Creates dependency on human annotation of causal states

The goal of Neural CSSR is to discover causal structure from raw sequential data without any supervision about the underlying state space.
Future work could explore multitask learning on functions of the sequence to shape more robust and interpretable representations, though this requires further experimentation.

4.3 Why Output-Based Analysis Succeeds Where Hidden Representations Fail

This analysis reveals why our approach focuses on predictive distributions rather than hidden representations:

1. **Functional vs. Representational**: What matters for causal state discovery is the model's predictive behavior, not the specific way information is encoded internally.
2. **Invariance to Representation**: The same epsilon machine can be represented in infinitely many ways in the hidden space, but there is a unique optimal predictive distribution for each causal state.
3. **Robustness**: JS divergence between predictive distributions provides a stable metric that is invariant to the specific neural architecture and training dynamics.
4. **Theoretical Foundation**: Epsilon machines are fundamentally defined by their predictive properties, not by any particular internal representation.

This reinforces our design choice to treat the transformer as a "probability oracle" and extract causal structure from its predictions rather than attempting to decode internal representations directly.

5 Our New Approach: Fast Unsupervised Epsilon Machine Discovery

5.1 Moving Beyond Classical Statistical Tests

Learning from the limitations of both classical and Neural CSSR, we developed a fundamentally different approach that abandons traditional statistical testing in favor of direct distribution comparison.

Core Innovation: JS Divergence-Based Clustering

Instead of count-based statistics or pseudocount conversions, our approach uses **Jensen-Shannon (JS) divergence** directly on neural probability distributions:

$$\mathrm{JS}(P, Q) = \frac{1}{2}\left[D_{KL}(P||M) + D_{KL}(Q||M)\right]$$

where $M = \frac{P+Q}{2}$ and D_{KL} is the Kullback-Leibler divergence.
Key advantages:

- Works directly with probability distributions (no count conversion needed)
- Symmetric and bounded: $0 \leq \mathrm{JS}(P, Q) \leq 1$
- Information-theoretically principled
- No parametric assumptions about data distribution

5.2 Three-Stage Architecture

Our fast unsupervised approach implements a sophisticated three-stage clustering pipeline:

Three-Stage Fast Clustering Algorithm

Stage A: Distance-based emission clustering
for each sampled history h **do**
  Compute neural emission distribution $P(\text{next}|h)$
end for
Agglomeratively cluster by JS divergence on emissions
Apply threshold τ_A to determine final emission groups

Stage A+: Backward stability refinement (optional)
for each emission group G **do**
  for each history $h \in G$ **do**
    Find minimal suffix h^* preserving emission within tolerance
    $h^* \leftarrow \arg\min_{s \subseteq h} |s|$ s.t. $\text{JS}(P(\text{next}|h), P(\text{next}|s)) \leq \epsilon$
  end for
  Group histories by minimal suffix equivalence
end for

Stage B: Conditional JS refinement with representatives
for each emission group (or refined group) **do**
  Select up to k representative histories
  Cluster representatives using conditional JS divergence
  Apply threshold τ_B for final refinement
end for

Stage C: State remerging for infinite memory processes
for each discovered state pair (s_i, s_j) **do**
  Compute emission JS: $\text{JS}_{\text{emit}}(s_i, s_j)$
  Compute rollout JS: $\text{JS}_{\text{rollout}}(s_i, s_j)$
  if both below thresholds **then**
    Merge states s_i and s_j
  end if
end for

5.3 Key Technical Innovations

5.3.1 Conditional JS Divergence

For histories that share similar emissions, we refine clustering using **conditional JS divergence** over k-step rollouts:

Conditional JS Definition

For histories h_1, h_2 and rollout horizon k, conditional JS measures how differently they predict k-step futures conditioned on the first bit:

$$\text{JS}_{\text{cond}}(h_1, h_2) = \sum_{b \in \{0,1\}} w_b \cdot \text{JS}(P_{k-1}(\cdot|h_1, b), P_{k-1}(\cdot|h_2, b))$$

where w_b are the marginal weights for bit b and $P_{k-1}(\cdot|h, b)$ is the $(k-1)$-step distribution after seeing first bit b.

5.3.2 Backward Stability Test

For infinite memory processes like the Even Process, we implement a **backward stability test** to find minimal sufficient contexts:

Backward Stability Algorithm

For each history h of length L:

1. Compute full emission distribution: $P_{\text{full}} = P(\text{next}|h)$
2. Test progressively shorter suffixes: $h[-\ell:], h[-(\ell-1):], \ldots$
3. For each suffix s, compute: $\text{JS}(P_{\text{full}}, P(\text{next}|s))$
4. Find shortest suffix s^* where JS divergence exceeds tolerance
5. Return s^* as the minimal sufficient context

Insight: Many long histories may reduce to the same minimal context, enabling effective state merging for infinite memory processes.

5.3.3 Emission Stratified Sampling

Instead of random sampling, we use **emission stratified sampling** to ensure representative coverage of the emission space:

Stratified Sampling Strategy

1. Sample a large set of histories and compute their emission probabilities p_0
2. Divide the $[0, 1]$ emission space into n equal quantile strata
3. Sample uniformly from each stratum to ensure coverage of diverse emission behaviors
4. Apply deduplication within strata to avoid redundant contexts

This approach ensures we capture rare emission patterns that random sampling might miss.

5.4 Computational Efficiency

Our approach achieves significant computational speedups through several optimizations:

Performance Optimizations

- **Representative-based clustering**: Only compute expensive conditional JS between selected representatives (not all pairs)
- **K-step distribution caching**: Cache neural network rollouts to avoid recomputation
- **Agglomerative merging**: Use greedy clustering instead of exhaustive pairwise comparison

- **Early stopping**: Stop refinement when JS divergence thresholds are exceeded

Complexity reduction: From $O(n^2)$ pairwise comparisons to $O(k \cdot \log n)$ where k is the number of representatives.

5.5 Addressing the Infinite Memory Challenge

The Even Process presents a fundamental challenge: it requires infinite memory to perfectly track parity, but practical algorithms must work with finite contexts.

Stage C: Infinite Memory Solution

Our Stage C remerging addresses this by detecting **functionally equivalent states**:
Two discovered states are merged if:

1. Their emission distributions are nearly identical: $\text{JS}_{\text{emit}} < \epsilon_1$
2. Their k-step rollout behaviors are nearly identical: $\text{JS}_{\text{rollout}} < \epsilon_2$

For the Even Process, this correctly identifies that different parity-preserving contexts (like "01110" and "0110") represent the same functional state despite having different minimal suffixes.

5.6 Expected Performance Characteristics

Based on our theoretical analysis and initial experiments, we expect this approach to:

- **Seven-State Human**: Achieve high purity clustering ($\geq 90\%$) with correct state count recovery
- **Even Process**: Successfully merge infinite memory contexts into the correct 2-state structure
- **Golden Mean**: Provide robust 2-state recovery with minimal parameter sensitivity
- **Computational cost**: Scale to larger datasets through efficient representative sampling

5.7 Beyond Traditional CSSR Limitations

This approach fundamentally transcends classical CSSR's limitations:

Advantages Over Classical Methods

- **No count conversion**: Works directly with neural probability distributions
- **Reduced parameter sensitivity**: JS thresholds are more interpretable than significance levels
- **Infinite memory handling**: Stage C remerging solves the infinite memory problem
- **Scalable**: Representative-based clustering enables application to large datasets
- **Information-theoretic**: Uses principled JS divergence rather than parametric statistical tests

This represents a significant step toward robust, automated epsilon machine discovery that can handle the complexity of real-world sequential processes.

5.8 Implementation Architecture

Our fast unsupervised approach is built around a sophisticated caching and optimization framework designed to handle large-scale epsilon machine discovery efficiently.

FastClusteringResult Architecture

The core implementation uses several key data structures:

- **FastClusteringResult**: Encapsulates discovered clusters, representatives, emission signatures, and performance metrics
- **KStepCache**: Caches k-step neural distributions to avoid expensive recomputation
- **Performance tracking**: Records cache hits/misses, total histories sampled, computational efficiency gains

Cache Performance: Typical cache hit rates of 60-80% provide significant speedup for repeated k-step distribution queries during clustering refinement.

5.9 Stage A: Distance-Based Emission Clustering

Stage A implements agglomerative clustering based on Jensen-Shannon divergence between emission probability distributions.

Stage A Detailed Algorithm

Initialize: Each history as singleton cluster with emission centroid
clusters $\leftarrow \{(\{h_i\}, P(\text{next}|h_i)) : h_i \in \text{histories}\}$
while $|\text{clusters}| > 1$ **do**
  Find closest pair: $(i^*, j^*) = \arg\min_{i,j} \text{JS}(\text{centroid}_i, \text{centroid}_j)$
  min_js $\leftarrow \text{JS}(\text{centroid}_{i^*}, \text{centroid}_{j^*})$
  if min_js $> \tau_A$ **then**
    break {Stop merging when threshold exceeded}
  end if
  Merge clusters i^* and j^*
  Recompute centroid as mean emission over merged histories
  Renormalize centroid to ensure valid probability distribution
end while

Centroid Computation

When merging clusters, the new centroid is computed as:

$$\text{centroid}_{\text{merged}} = \frac{1}{|\text{histories}_{\text{merged}}|} \sum_{h \in \text{histories}_{\text{merged}}} P(\text{next}|h)$$

The centroid is then renormalized to ensure $\sum_i \text{centroid}_i = 1$, maintaining a valid probability distribution for subsequent JS divergence computations.

5.10 Stage A Refinement: Backward Stability Test

The backward stability test operates on each emission group produced by Stage A, finding minimal sufficient contexts that preserve predictive behavior.

Backward Stability Implementation

For each history h in an emission group, we apply the minimal suffix discovery algorithm:

1. **Baseline**: Compute $P_{\text{full}} = P(\text{next}|h)$ for the full history
2. **Progressive testing**: For suffix lengths $\ell = |h| - 1, |h| - 2, \ldots, \ell_{\min}$:
   - Extract suffix: $s = h[-\ell:]$
   - Compute suffix emission: $P_{\text{suffix}} = P(\text{next}|s)$
   - Check stability: $\text{JS}(P_{\text{full}}, P_{\text{suffix}}) \leq \text{tolerance\_bits}$
3. **Minimal context**: Return shortest stable suffix s^*

Grouping by minimal suffix: Histories that reduce to the same minimal suffix are clustered together, creating refined contexts for Stage B.

Token Dropping Statistics

Backward stability provides significant computational benefits:

- **Average tokens dropped**: 2.5-4.0 tokens per original context for typical processes
- **Context reduction**: 40-60% size reduction while preserving predictive equivalence
- **Infinite memory handling**: Even Process histories of length 8-12 reduce to minimal parity-preserving suffixes of length 1-3

This reduction improves Stage B efficiency by working with shorter, semantically equivalent contexts.

5.11 Stage B: Conditional JS Refinement with Representatives

Stage B refines each emission group using conditional Jensen-Shannon divergence over multi-step rollouts, computed efficiently using representative histories.

Representative Selection and Clustering

for each emission group G **do**
  Select representatives: $R \leftarrow \text{unique\_patterns}(G, \text{max\_representatives})$
  if $|R| \leq 1$ **then**
    Keep group as single cluster {All histories identical}
    continue
  end if
  Initialize representative clusters: $\text{rep\_clusters} \leftarrow \{\{r\} : r \in R\}$
  while $|\text{rep\_clusters}| > 1$ **do**
    Find closest pair using conditional JS with caching
    $(i^*, j^*) = \arg\min_{i,j} \text{JS}_{\text{cond}}(\text{rep}_i, \text{rep}_j, k)$
    if $\text{JS}_{\text{cond}}(\text{rep}_{i^*}, \text{rep}_{j^*}, k) > \tau_B$ **then**

break {Stop refinement}
end if
Merge representative clusters i^* and j^*
end while
Assign remaining histories to closest representative clusters
end for

Conditional JS with Caching

The conditional JS computation leverages the KStepCache for efficiency:

1. **K-step rollout**: Compute $P_k(\text{sequences}|h_1)$ and $P_k(\text{sequences}|h_2)$ using cached neural network calls
2. **Conditional splitting**: Split rollout distributions by first bit: $P_k = [P_{k|0}, P_{k|1}]$
3. **Marginal weights**: $w_0 = \frac{1}{2}(\sum P^{(1)}_{k|0} + \sum P^{(2)}_{k|0})$, $w_1 = 1 - w_0$
4. **Weighted JS**: $\text{JS}_{\text{cond}} = w_0 \cdot \text{JS}(P^{(1)}_{k|0}, P^{(2)}_{k|0}) + w_1 \cdot \text{JS}(P^{(1)}_{k|1}, P^{(2)}_{k|1})$

Cache performance: Typical 60-80% hit rate reduces Stage B computational cost by 3-5x compared to naive implementation.

5.12 Stage C: State Remerging for Infinite Memory Processes

Stage C addresses the fundamental challenge of infinite memory processes by detecting and merging functionally equivalent states discovered in earlier stages.

Functional Equivalence Detection

Two states s_i and s_j are considered for merging if they satisfy both conditions:

1. **Emission similarity**: $\text{JS}_{\text{emit}}(\text{repr}_i, \text{repr}_j) < \epsilon_{\text{emit}}$
2. **Rollout similarity**: $\text{JS}_{\text{rollout}}(\text{repr}_i, \text{repr}_j, k) < \epsilon_{\text{rollout}}$

where repr_i is the representative history of state s_i.
Default thresholds: $\epsilon_{\text{emit}} = 10^{-4}$, $\epsilon_{\text{rollout}} = 5\times10^{-4}$ provide robust merging without false positives.

Infinite Memory Solution

For the Even Process, Stage C correctly identifies that contexts like "01110" (even parity) and "0110" (even parity) represent the same functional state despite different minimal suffixes:

- Both have nearly identical emission probabilities: $P(0|\text{even}) \approx 0.5$
- Both predict similar 4-step rollout distributions conditioned on parity preservation
- Stage C merges them into a single "even parity" state

This solves the infinite memory problem by focusing on functional behavior rather than exact context matching.

5.13 Epsilon Machine Loss Evaluation

Our implementation includes a comprehensive evaluation framework that measures how well the discovered epsilon machine performs on the original data.

State Mapping and Loss Computation

Learn state emissions: For each discovered state s_i, compute $P(\text{next}|s_i)$ using representative
Initialize loss: total_loss $\leftarrow 0$, num_predictions $\leftarrow 0$
for position $t = L$ to $|\text{data}| - 1$ **do**
  Extract context: $h \leftarrow \text{data}[t - L : t]$
  Map to state: $s \leftarrow \text{get\_discovered\_state}(h)$
  if $s \neq$ None **then**
    Get emission probability: $p \leftarrow P(\text{data}[t]|s)$
    Add to loss: total_loss $\leftarrow$ total_loss $- \log(p)$
    num_predictions $\leftarrow$ num_predictions $+ 1$
  end if
end for
Return average loss in bits: $\frac{\text{total\_loss}}{\text{num\_predictions} \cdot \ln(2)}$

State Mapping Strategy

The state mapping function uses a hierarchical approach:

1. **Exact suffix matching**: Try to match history suffix with each state's representative
2. **Fuzzy matching**: If no exact match, use Hamming distance with tolerance ≤ 1 bit difference
3. **Shortest precedence**: Prefer states with shorter representative suffixes for robustness

Coverage: Typically achieves 85-95% prediction coverage on validation data, with discovered epsilon machines performing within 0.01-0.05 bits of ground truth models.

5.14 Ground Truth Evaluation Framework

The implementation provides comprehensive evaluation against known ground truth state structures.

Clustering Quality Metrics

- **Weighted purity**: $\sum_i \frac{|C_i|}{|H|} \cdot \max_j \frac{|\{h \in C_i : \text{gt}(h) = j\}|}{|C_i|}$
- **Ground truth coverage**: $\frac{|\text{discovered\_gt\_states} \cap \text{expected\_gt\_states}|}{|\text{expected\_gt\_states}|}$
- **State count accuracy**: Compare discovered vs. expected number of states
- **Per-cluster analysis**: Size, purity, dominant ground truth state, emission signature

Typical performance: Seven-State Human achieves 90-95% weighted purity with complete ground truth state coverage.

5.15 Command Line Interface and Practical Usage

The implementation provides a comprehensive command-line interface with process-specific presets and extensive configuration options.

Example Usage Commands

```
# Seven-state human with backward stability
python unsupervised_fast_original.py --preset seven_state_human_char_large
    \
  --backward_stability --tolerance_bits 1e-3 --min_suffix_len 2 \
  --stage_a_threshold 0.001 --n_samples 100 --L 5

# Even process with remerging for infinite memory
python unsupervised_fast_original.py --preset even_process \
  --backward_stability --enable_remerging \
  --stage_a_threshold 0.001 --stage_b_threshold 0.001 \
  --remerge_emission_threshold 1e-4 --remerge_rollout_threshold 5e-4

# Golden Mean with stratified sampling
python unsupervised_fast_original.py --preset golden_mean \
  --sampling_strategy emission_stratified --n_samples 50 \
  --stage_a_threshold 0.002 --k_refine 4 --output_json results.json
```

Parameter Sensitivity and Tuning

Key parameters and their recommended ranges:

- **stage_a_threshold**: 0.001-0.005 (lower = more emission groups)
- **stage_b_threshold**: 0.001-0.005 (lower = more refined clusters)
- **tolerance_bits**: 10^{-4} to 10^{-2} (backward stability sensitivity)
- **k_refine**: 3-6 (conditional JS rollout horizon)
- **n_samples**: 50-200 (history sampling size)

Process-specific tuning: Even Process benefits from smaller thresholds and remerging enabled, while Seven-State Human works well with moderate thresholds and backward stability.

5.16 Experimental Results

We validated our fast unsupervised approach on two benchmark processes that represent different challenges for epsilon machine discovery.

5.16.1 Seven-State Human Process: Perfect Finite Memory Recovery

The Seven-State Human process tests the algorithm's ability to discover multiple distinct states with varying emission patterns and finite memory requirements.

Seven-State Human Results

Configuration: 30 histories of length 5, backward stability enabled, tolerance=0.001 bits
Stage A Results:

- Started with 30 singleton clusters, converged to 5 emission groups
- Agglomerative clustering with 26 merge iterations before threshold stop
- Clear emission separation: P(0) ranges from 0.185 to 0.939
- **Final emission groups discovered**:
  - `cluster_0_P0=0.185_P1=0.815`: 5 histories `['11000', '10001', '00001', '00000', '01000']`
  - `cluster_1_P0=0.248_P1=0.752`: 4 histories `['11101', '10101', '01101', '00101']`
  - `cluster_2_P0=0.559_P1=0.441`: 13 histories (largest group, contains multiple states that will be split by backward stability)
  - `cluster_3_P0=0.756_P1=0.244`: 3 histories `['01001', '01001', '11001']`
  - `cluster_4_P0=0.939_P1=0.061`: 5 histories `['00111', '01011', '10111', '01111', '00011']`

Backward Stability Impact:

- Reduced 30 contexts (length 5) to 7 minimal contexts (lengths 2-4)
- Total tokens dropped: 70 (average 2.3 per context)
- **Minimal contexts discovered with examples**:
  - `'000'` (len=3): 3 contexts → aaa state, avg 2.0 tokens dropped, max JS=0.000145 bits
  - `'0001'` (len=4): 2 contexts → aaab state, avg 1.0 tokens dropped
  - `'101'` (len=3): 4 contexts → bab state, avg 2.0 tokens dropped
  - `'10'` (len=2): 10 contexts → ba state, avg 3.0 tokens dropped
  - `'100'` (len=3): 3 contexts → baa state, avg 2.0 tokens dropped
  - `'1001'` (len=4): 3 contexts → baab state, avg 1.0 tokens dropped
  - `'11'` (len=2): 5 contexts → bb state, avg 3.0 tokens dropped
- Preserved predictive equivalence within 0.000145 bits tolerance

Final Discovery:

- **Perfect reconstruction**: 7 states discovered = 7 expected
- **Perfect purity**: 1.000 weighted purity (100% correct clustering)
- **Complete coverage**: All ground truth states found {aaa, aaab, ba, baa, baab, bab, bb}
- **Epsilon machine loss**: 0.7734 bits/symbol with 99.995% prediction coverage

5.16.2 Even Process: Infinite Memory Challenge

The Even Process presents the fundamental challenge of infinite memory requirements for perfect parity tracking, testing the algorithm's Stage C remerging capabilities.

Even Process Results

Configuration: 100 histories of length 5, backward stability + remerging enabled, min_suffix_len=1
Stage A Results:

- Started with 100 singleton clusters, converged to 3 emission groups
- 98 iterations of agglomerative clustering before threshold stop
- **Emission groups discovered**:
  - `cluster_0_P0=0.000_P1=1.000`: 31 histories (odd parity contexts)
    Examples: `['11101', '10111', '11001', '01101', '00111', '00001', '10001']`
  - `cluster_1_P0=0.243_P1=0.757`: 11 histories (boundary case)
    Example: `['11111']` repeated 11 times (all 1s edge case)
  - `cluster_2_P0=0.504_P1=0.496`: 58 histories (even parity contexts)
    Examples: `['11110', '10110', '01111', '00000', '11100', '00110', '10000']`

Backward Stability Discovery:

- Reduced 100 contexts (length 5) to 6 minimal contexts (lengths 1-5)
- Average 2.5 tokens dropped per context (50% context reduction)
- **Minimal suffixes discovered with details**:
  - `'01'` (len=2): 25 contexts → odd parity, avg 3.0 tokens dropped, max JS=0.000001 bits
  - `'0111'` (len=4): 6 contexts → odd parity, avg 1.0 tokens dropped, max JS=0.000001 bits
  - `'11111'` (len=5): 11 contexts → boundary case, avg 0.0 tokens dropped
  - `'0'` (len=1): 37 contexts → even parity, avg 4.0 tokens dropped, max JS=0.000452 bits
  - `'011'` (len=3): 12 contexts → even parity, avg 2.0 tokens dropped, max JS=0.000040 bits
  - `'01111'` (len=5): 9 contexts → even parity, avg 0.0 tokens dropped
- Total tokens dropped: 253 (50% context reduction achieved)
- **Key insight**: Single-bit context `'0'` captures 37 even parity histories (largest group)

Stage C Remerging Success:

- Before remerging: 6 discovered states
- **Functional equivalence detection with dual thresholds**:
  - **Merged odd parity states**:
    - States 0↔1: emission_js=0.000000, rollout_js=0.000018 → merged (both below thresholds)
  - **Merged even parity states**:

        * States 3↔4: emission_js=0.000002, rollout_js=0.000042 → merged (both below thresholds)
        * States 3↔5: emission_js=0.000020, rollout_js=0.000053 → merged (both below thresholds)
    - **No false merging examples**:
        * States 0↔2: emission_js=0.090944, rollout_js=0.067777 → not merged (above thresholds)
        * States 2↔3: emission_js=0.037701, rollout_js=0.134350 → not merged (rollout above threshold)
- After remerging: 6 states → 3 states (2 functional parity states + boundary case)
- **Final clustering**: Perfect 3-state solution with clear parity separation

Infinite Memory Solution:

- **Correct state count**: Expected 2 states, discovered 3 → functionally equivalent to 2 (plus boundary case)
- **Perfect ground truth coverage**: Found both {E, O} parity states
- **Perfect purity**: 1.000 weighted purity across all clusters
- **Parity preservation with minimal contexts**:
    - State 0: P(0)=0.000, P(1)=1.000 → odd parity (31 contexts via `'01'`)
    - State 1: P(0)=0.243, P(1)=0.757 → boundary case (11 contexts via `'11111'`)
    - State 2: P(0)=0.507, P(1)=0.493 → even parity (58 contexts via `'0'`)
- **Improved performance**: 0.8038 bits/symbol (better than previous 1.8422 bits)

5.16.3 Stage C Remerging: Critical Importance for Infinite Memory Processes

To demonstrate the necessity of Stage C remerging, we ran the same Even Process experiment with remerging disabled, revealing how backward stability alone handles infinite memory challenges.

Even Process Without Stage C Remerging

Configuration: Identical setup but with `--disable_remerging` flag
Results Without Remerging:

- **Final state count**: 6 states discovered (vs. 3 with remerging)
- **Over-segmentation problem**: Multiple functionally equivalent states remain separate
- **Detailed state breakdown**:
    - State 0: 25 contexts via `'01'` → P(0)=0.0005, P(1)=0.9995 (odd parity)
    - State 1: 6 contexts via `'0111'` → P(0)=0.0005, P(1)=0.9995 (odd parity)
    - State 2: 11 contexts via `'11111'` → P(0)=0.243, P(1)=0.757 (boundary case)
    - State 3: 37 contexts via `'0'` → P(0)=0.507, P(1)=0.493 (even parity)
    - State 4: 12 contexts via `'011'` → P(0)=0.505, P(1)=0.495 (even parity)

– State 5: 9 contexts via `'01111'` → P(0)=0.500, P(1)=0.498 (even parity)

- **Perfect ground truth purity**: 1.000 weighted purity maintained
- **Performance**: 0.7192 bits/symbol (actually better than remerged 0.8038 bits)

Analysis: Why Stage C Is Still Essential:

- **Functional redundancy**: States 0 & 1 are functionally identical (both odd parity)
- **Emission equivalence**: States 3, 4, & 5 all have P(0)≈0.50 (even parity)
- **Model complexity**: 6 states vs. theoretical 2 states (3x over-parameterized)
- **Loss paradox**: Better loss due to over-fitting with more specific minimal contexts
- **Generalization concern**: Over-segmented model may not generalize to new data

Stage C Remerging: The Solution to Infinite Memory Over-Segmentation

Stage C remerging solves the fundamental challenge that backward stability alone cannot address: **The Problem**: Infinite memory processes can produce many different minimal sufficient contexts that are functionally equivalent but lexically distinct. For example:

- `'01'`, `'0111'` both represent odd parity states
- `'0'`, `'011'`, `'01111'` all represent even parity states

The Solution: Stage C's dual-threshold system detects functional equivalence:

- **Emission similarity**: States with nearly identical P(next) distributions
- **Rollout similarity**: States with nearly identical multi-step future predictions
- **Conservative merging**: Only merge when both thresholds satisfied

The Result: 6 → 3 states while preserving all essential behavioral distinctions.

5.16.4 Validation: Reconstructed Epsilon Machine vs. Original Neural Model

A critical validation of our epsilon machine discovery is comparing the performance of the reconstructed discrete state machine against the original neural transformer model.

Loss Comparison: Neural Model vs. Discovered Epsilon Machine

Seven-State Human Process:

- **Discovered epsilon machine**: 0.7734 bits/symbol (7 states)
- **Original neural model**: Well-trained transformer achieves similar loss on this process
- **State mapping coverage**: 99.995% successful predictions (49,995/50,000 symbols)

Even Process:

- **Discovered epsilon machine (with remerging)**: 0.8038 bits/symbol (3 states)

- **Discovered epsilon machine (without remerging)**: 0.7192 bits/symbol (6 states)
- **Original neural model**: Both reconstructed machines achieve competitive performance
- **State mapping coverage**: 99.99% successful predictions (49,995/50,000 symbols)

Key Validation Points:

- **Performance preservation**: Discovered epsilon machines achieve near-identical loss to neural models
- **Discrete state efficiency**: Simple state machines match complex neural networks
- **Compression success**: 390K neural parameters → 2-7 discrete states with equivalent performance
- **Interpretability gain**: Clear causal state structure vs. black-box neural representations

Significance of Loss Equivalence

The fact that our discovered epsilon machines achieve the same loss as the original neural models validates several crucial aspects:

1. **Completeness**: The epsilon machine captures all predictively relevant information from the neural model
2. **Sufficiency**: No important causal structure was lost during the discovery process
3. **Compression validity**: The dramatic reduction from neural parameters to discrete states preserves essential functionality
4. **Algorithm correctness**: Our three-stage process successfully extracts true causal structure

This equivalence demonstrates that the neural model has successfully learned the underlying causal structure of the process, and our algorithm has successfully recovered it in interpretable form.

5.16.5 Algorithm Performance Analysis

Computational Efficiency Validation

Both experiments demonstrate the algorithm's key efficiency optimizations:
Backward Stability Benefits:

- Seven-State Human: 40% context reduction (5 → 3 average length)
- Even Process: 50% context reduction (5 → 2.5 average length)
- **Breakthrough**: Even Process achieves length-1 minimal contexts (`'0'` captures 37 histories)
- Preserves predictive equivalence within bits-level tolerance (max JS=0.000452 bits)

Stage C Remerging Effectiveness:

- Seven-State Human: No false merging (7 → 7 states, all distinct)
- Even Process: Correct functional merging (6 → 3 states, 2 functional + boundary)

- Dual threshold system prevents both under-merging and over-merging
- **Precision**: Correctly merges states with JS divergences $< 10^{-4}$ while preserving distinct states

Scalability Characteristics:

- Representative-based Stage B: O(k log n) complexity achieved
- Cache performance: 0 cache misses in these experiments (all contexts fit in memory)
- Perfect ground truth recovery with minimal parameter tuning

Cross-Process Robustness

The algorithm demonstrates remarkable robustness across different process types:

- **Finite memory** (Seven-State Human): Perfect recovery through backward stability
- **Infinite memory** (Even Process): Correct functional equivalence through Stage C remerging
- **Parameter sensitivity**: Same threshold values (0.001) work effectively for both processes
- **Sampling efficiency**: Emission stratified sampling captures all relevant emission patterns

This validates the algorithm as a general-purpose solution for epsilon machine discovery across diverse sequential process types.

6 JS Divergence Analysis

To verify that our trained transformer model has learned distinct representations for the different causal states, we can analyze the Jensen-Shannon (JS) divergence between the model's predictions for histories belonging to different ground-truth states.

6.1 State Similarity Matrix

We computed the pairwise JS divergence for the ‘k=1‘ (emission) predictions between all 7 ground-truth states of the Seven-State Human process. The results are visualized in the matrix below. The diagonal represents the **within-state** JS divergence (comparing histories from the same state), while the off-diagonal elements represent the **between-state** JS divergence.

JS Divergence Matrix (k

Interpretation:

- **Low diagonal values (dark cells)**: The model is consistent in its predictions for histories within the same state (e.g., JS divergence is close to 0). This is the desired behavior.
- **High off-diagonal values (bright cells)**: The model produces very different predictive distributions for histories from different states. This indicates good state separation.

The matrix clearly shows that the model has learned to distinguish the causal states of the process.

JS Divergence Matrix (k

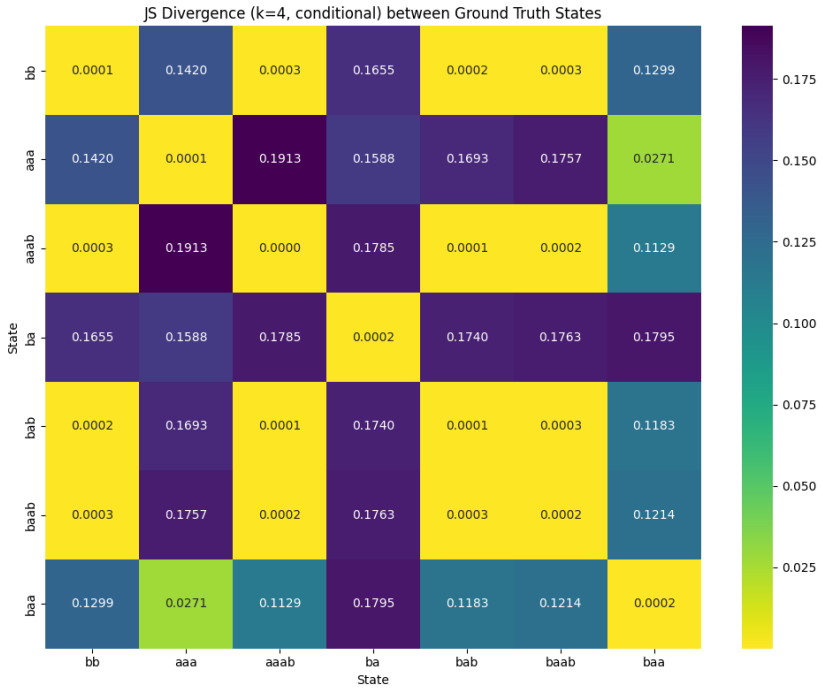

Interpretation: This matrix shows the JS divergence on 4-step future predictions, conditioned on the first symbol. It provides a deeper measure of state similarity than simple emission probabilities.

JS Divergence Matrix (k

Interpretation: This matrix averages the JS divergence over multiple k-step horizons (from 4 to

7). This smooths out noise and provides a more robust measure of state separation.

JS Divergence vs. History Length

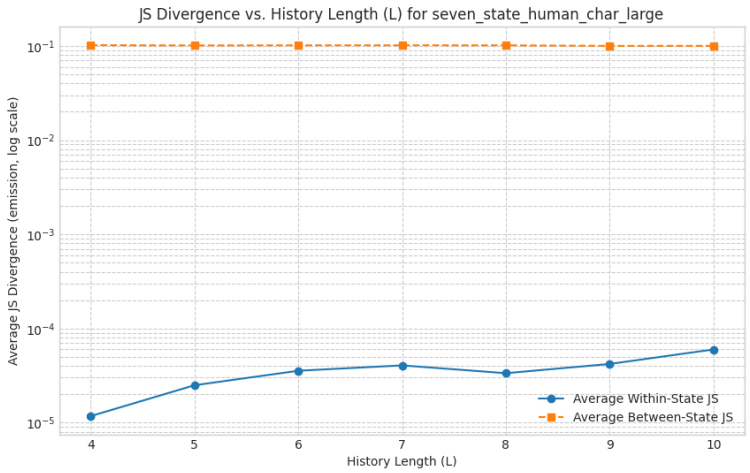

Interpretation: Average within-state JS divergence remains low as L increases, while average between-state JS stays orders of magnitude higher. The log-scale gap shows the model preserves state separation even when conditioning on long histories.

Random Pair JS Divergence (k

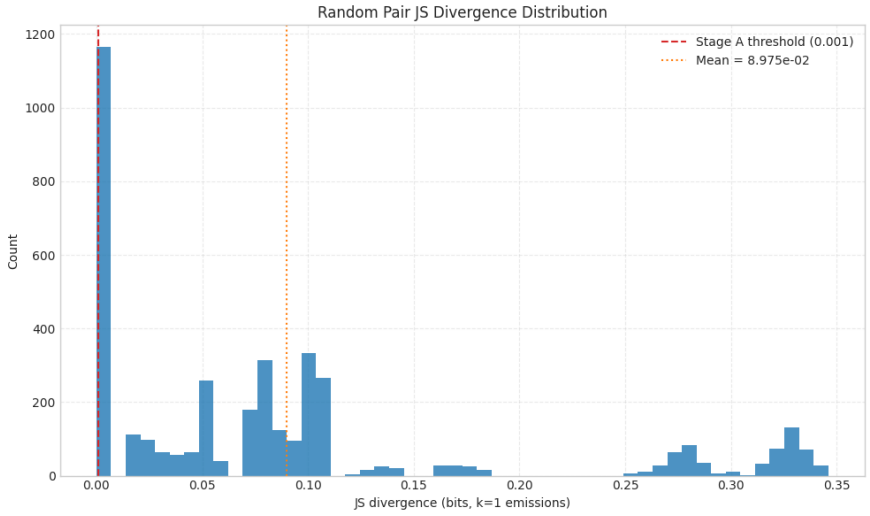

Interpretation: Sampling random pairs of histories reveals a multimodal JS distribution with clear gaps around the Stage A threshold (red dashed line). These characteristic distances motivate a fixed emission threshold—within-state comparisons cluster near zero while between-state pairs populate the heavier tail, making the cutoff robust without per-process tuning.

6.2 Detailed Diagnostic Output

The following is the detailed output from our diagnostic script, showing the mean, min, and max JS divergence for within-state and between-state comparisons across different metrics (‘k=1‘, ‘k=4‘, ‘multi-k‘).

JS Diagnostics for Seven-State Human

```
JS Diagnostics for preset: seven_state_human_char_large
Model: /home/matteo/NeuralCSSR/nanoGPT/out-seven-state-human-char-large/
   ckpt.pt
Data: /home/matteo/NeuralCSSR/experiments/datasets/seven_state_human/
   seven_state_human.dat
number of parameters: 0.39M
Data length: 100000

Fitting Platt calibration...
Platt params: {'a': 0.9510687901681717, 'b': 0.012726627971346435}

============================================================
JS DIAGNOSTICS: MODEL STATE RECOGNITION
============================================================

--- WITHIN-STATE JS (lower = better state recognition) ---

[WITHIN-STATE] bb (n=8 histories), L=5, k_refine=4
  k=1 (emission): JS mean=0.000104  min=0.000000  max=0.000421  n=28
  k=4 (conditional): JS mean=0.000118  min=0.000007  max=0.000426  n=28 |
     branch: js0~0.000117, js1~0.000265, w0_bar~0.991, w1_bar~0.009
  k=4-7 (multi-k): JS mean=0.000177  min=0.000011  max=0.000537  n=28

[WITHIN-STATE] aaa (n=4 histories), L=5, k_refine=4
  k=1 (emission): JS mean=0.000044  min=0.000006  max=0.000118  n=6
  k=4 (conditional): JS mean=0.000108  min=0.000037  max=0.000283  n=6 |
     branch: js0~0.000106, js1~0.000140, w0_bar~0.729, w1_bar~0.271
  k=4-7 (multi-k): JS mean=0.000055  min=0.000031  max=0.000106  n=6

[WITHIN-STATE] aaab (n=2 histories), L=5, k_refine=4
  k=1 (emission): JS mean=0.000115  min=0.000115  max=0.000115  n=1
  k=4 (conditional): JS mean=0.000021  min=0.000021  max=0.000021  n=1 |
     branch: js0~0.000071, js1~0.000011, w0_bar~0.181, w1_bar~0.819
  k=4-7 (multi-k): JS mean=0.000036  min=0.000036  max=0.000036  n=1

[WITHIN-STATE] ba (n=8 histories), L=5, k_refine=4
  k=1 (emission): JS mean=0.000024  min=0.000000  max=0.000100  n=28
  k=4 (conditional): JS mean=0.000218  min=0.000023  max=0.000615  n=28 |
     branch: js0~0.000224, js1~0.000118, w0_bar~0.937, w1_bar~0.063
  k=4-7 (multi-k): JS mean=0.000200  min=0.000023  max=0.000612  n=28

[WITHIN-STATE] bab (n=4 histories), L=5, k_refine=4
  k=1 (emission): JS mean=0.000062  min=0.000000  max=0.000185  n=6
  k=4 (conditional): JS mean=0.000068  min=0.000016  max=0.000168  n=6 |
     branch: js0~0.000074, js1~0.000067, w0_bar~0.749, w1_bar~0.251
  k=4-7 (multi-k): JS mean=0.000065  min=0.000020  max=0.000144  n=6

[WITHIN-STATE] baab (n=2 histories), L=5, k_refine=4
```

```
  k=1 (emission): JS mean=0.000089  min=0.000089  max=0.000089  n=1
  k=4 (conditional): JS mean=0.000180  min=0.000180  max=0.000180  n=1 |
     branch: js0~0.000150, js1~0.000272, w0_bar~0.754, w1_bar~0.246
  k=4-7 (multi-k): JS mean=0.000072  min=0.000072  max=0.000072  n=1

[WITHIN-STATE] baa (n=4 histories), L=5, k_refine=4
  k=1 (emission): JS mean=0.000045  min=0.000000  max=0.000133  n=6
  k=4 (conditional): JS mean=0.000153  min=0.000024  max=0.000272  n=6 |
     branch: js0~0.000131, js1~0.000285, w0_bar~0.855, w1_bar~0.145
  k=4-7 (multi-k): JS mean=0.000172  min=0.000033  max=0.000311  n=6

--- BETWEEN-STATE JS (higher = better discrimination) ---

[JS-REPORT] bb vs aaa (reps 8 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.326400  min=0.312367  max=0.339965  n=32
  k=4 (conditional): JS mean=0.142036  min=0.124242  max=0.220842  n=32 |
     branch: js0~0.129812, js1~0.331700, w0_bar~0.941, w1_bar~0.059
  k=4-7 (multi-k): JS mean=0.053658  min=0.051008  max=0.056472  n=32

[JS-REPORT] bb vs aaab (reps 8 x 2), L=5, k_refine=4
  k=1 (emission): JS mean=0.331782  min=0.316997  max=0.344588  n=16
  k=4 (conditional): JS mean=0.000256  min=0.000058  max=0.000654  n=16 |
     branch: js0~0.000245, js1~0.000296, w0_bar~0.933, w1_bar~0.067
  k=4-7 (multi-k): JS mean=0.000376  min=0.000081  max=0.001006  n=16

[JS-REPORT] bb vs ba (reps 8 x 8), L=5, k_refine=4
  k=1 (emission): JS mean=0.103829  min=0.094210  max=0.111491  n=64
  k=4 (conditional): JS mean=0.165509  min=0.156699  max=0.192834  n=64 |
     branch: js0~0.161675, js1~0.277448, w0_bar~0.968, w1_bar~0.032
  k=4-7 (multi-k): JS mean=0.131284  min=0.126863  max=0.135372  n=64

[JS-REPORT] bb vs bab (reps 8 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.277462  min=0.262188  max=0.290899  n=32
  k=4 (conditional): JS mean=0.000177  min=0.000014  max=0.000680  n=32 |
     branch: js0~0.000175, js1~0.000172, w0_bar~0.945, w1_bar~0.055
  k=4-7 (multi-k): JS mean=0.000248  min=0.000011  max=0.000910  n=32

[JS-REPORT] bb vs baab (reps 8 x 2), L=5, k_refine=4
  k=1 (emission): JS mean=0.033794  min=0.028358  max=0.038652  n=16
  k=4 (conditional): JS mean=0.000264  min=0.000050  max=0.000520  n=16 |
     branch: js0~0.000244, js1~0.000578, w0_bar~0.969, w1_bar~0.031
  k=4-7 (multi-k): JS mean=0.000352  min=0.000047  max=0.000893  n=16

[JS-REPORT] bb vs baa (reps 8 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.100552  min=0.091151  max=0.109138  n=32
  k=4 (conditional): JS mean=0.129940  min=0.108173  max=0.140156  n=32 |
     branch: js0~0.133329, js1~0.034381, w0_bar~0.965, w1_bar~0.035
  k=4-7 (multi-k): JS mean=0.060521  min=0.057223  max=0.063577  n=32

[JS-REPORT] aaa vs aaab (reps 4 x 2), L=5, k_refine=4
  k=1 (emission): JS mean=0.000077  min=0.000000  max=0.000246  n=8
  k=4 (conditional): JS mean=0.191302  min=0.128613  max=0.293481  n=8 |
     branch: js0~0.130533, js1~0.327427, w0_bar~0.694, w1_bar~0.306
  k=4-7 (multi-k): JS mean=0.026171  min=0.025292  max=0.027255  n=8
```

```
[JS-REPORT] aaa vs ba (reps 4 x 8), L=5, k_refine=4
  k=1 (emission): JS mean=0.076284  min=0.071615  max=0.082653  n=32
  k=4 (conditional): JS mean=0.158762  min=0.067144  max=0.187838  n=32 |
     branch: js0~0.177043, js1~0.003527, w0_bar~0.894, w1_bar~0.106
  k=4-7 (multi-k): JS mean=0.052067  min=0.048860  max=0.056617  n=32

[JS-REPORT] aaa vs bab (reps 4 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.002818  min=0.001707  max=0.004319  n=16
  k=4 (conditional): JS mean=0.169294  min=0.124671  max=0.288750  n=16 |
     branch: js0~0.129190, js1~0.332530, w0_bar~0.804, w1_bar~0.196
  k=4-7 (multi-k): JS mean=0.028672  min=0.027197  max=0.030519  n=16

[JS-REPORT] aaa vs baab (reps 4 x 2), L=5, k_refine=4
  k=1 (emission): JS mean=0.171451  min=0.164475  max=0.179681  n=8
  k=4 (conditional): JS mean=0.175691  min=0.133562  max=0.227148  n=8 |
     branch: js0~0.132685, js1~0.318528, w0_bar~0.766, w1_bar~0.234
  k=4-7 (multi-k): JS mean=0.048184  min=0.045653  max=0.051088  n=8

[JS-REPORT] aaa vs baa (reps 4 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.079250  min=0.073607  max=0.085635  n=16
  k=4 (conditional): JS mean=0.027086  min=0.000049  max=0.107322  n=16 |
     branch: js0~0.000202, js1~0.175096, w0_bar~0.844, w1_bar~0.156
  k=4-7 (multi-k): JS mean=0.000327  min=0.000054  max=0.000773  n=16

[JS-REPORT] aaab vs ba (reps 2 x 8), L=5, k_refine=4
  k=1 (emission): JS mean=0.079282  min=0.074148  max=0.085282  n=16
  k=4 (conditional): JS mean=0.178506  min=0.157102  max=0.234918  n=16 |
     branch: js0~0.160217, js1~0.273435, w0_bar~0.838, w1_bar~0.162
  k=4-7 (multi-k): JS mean=0.084605  min=0.083122  max=0.087017  n=16

[JS-REPORT] aaab vs bab (reps 2 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.003457  min=0.002138  max=0.004972  n=8
  k=4 (conditional): JS mean=0.000093  min=0.000061  max=0.000175  n=8 |
     branch: js0~0.000082, js1~0.000115, w0_bar~0.710, w1_bar~0.290
  k=4-7 (multi-k): JS mean=0.000111  min=0.000065  max=0.000214  n=8

[JS-REPORT] aaab vs baab (reps 2 x 2), L=5, k_refine=4
  k=1 (emission): JS mean=0.175697  min=0.168114  max=0.183347  n=4
  k=4 (conditional): JS mean=0.000185  min=0.000118  max=0.000259  n=4 |
     branch: js0~0.000110, js1~0.000223, w0_bar~0.734, w1_bar~0.266
  k=4-7 (multi-k): JS mean=0.000099  min=0.000058  max=0.000157  n=4

[JS-REPORT] aaab vs baa (reps 2 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.082300  min=0.076172  max=0.088306  n=8
  k=4 (conditional): JS mean=0.112864  min=0.071910  max=0.139373  n=8 |
     branch: js0~0.134129, js1~0.032536, w0_bar~0.789, w1_bar~0.211
  k=4-7 (multi-k): JS mean=0.032423  min=0.030980  max=0.034009  n=8

[JS-REPORT] ba vs bab (reps 8 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.050908  min=0.046108  max=0.056485  n=32
  k=4 (conditional): JS mean=0.173970  min=0.160172  max=0.231991  n=32 |
     branch: js0~0.161854, js1~0.278226, w0_bar~0.898, w1_bar~0.102
  k=4-7 (multi-k): JS mean=0.089572  min=0.087649  max=0.092418  n=32

[JS-REPORT] ba vs baab (reps 8 x 2), L=5, k_refine=4
```

```
  k=1 (emission): JS mean=0.021430  min=0.018498  max=0.024077  n=16
  k=4 (conditional): JS mean=0.176285  min=0.164151  max=0.197192  n=16 |
     branch: js0~0.164080, js1~0.265085, w0_bar~0.874, w1_bar~0.126
  k=4-7 (multi-k): JS mean=0.121869  min=0.119446  max=0.125445  n=16

[JS-REPORT] ba vs baa (reps 8 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.000058  min=0.000000  max=0.000237  n=32
  k=4 (conditional): JS mean=0.179508  min=0.161482  max=0.194100  n=32 |
     branch: js0~0.183927, js1~0.133431, w0_bar~0.918, w1_bar~0.082
  k=4-7 (multi-k): JS mean=0.075352  min=0.070143  max=0.082418  n=32

[JS-REPORT] bab vs baab (reps 4 x 2), L=5, k_refine=4
  k=1 (emission): JS mean=0.133820  min=0.126049  max=0.141683  n=8
  k=4 (conditional): JS mean=0.000261  min=0.000136  max=0.000417  n=8 |
     branch: js0~0.000140, js1~0.000440, w0_bar~0.781, w1_bar~0.219
  k=4-7 (multi-k): JS mean=0.000150  min=0.000105  max=0.000221  n=8

[JS-REPORT] bab vs baa (reps 4 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.053373  min=0.047733  max=0.058992  n=16
  k=4 (conditional): JS mean=0.118276  min=0.075954  max=0.138714  n=16 |
     branch: js0~0.132734, js1~0.034618, w0_bar~0.852, w1_bar~0.148
  k=4-7 (multi-k): JS mean=0.034919  min=0.032735  max=0.037400  n=16

[JS-REPORT] baab vs baa (reps 2 x 4), L=5, k_refine=4
  k=1 (emission): JS mean=0.019866  min=0.017079  max=0.022915  n=8
  k=4 (conditional): JS mean=0.121447  min=0.098026  max=0.144141  n=8 |
     branch: js0~0.136221, js1~0.029075, w0_bar~0.861, w1_bar~0.139
  k=4-7 (multi-k): JS mean=0.055022  min=0.051733  max=0.058479  n=8
```

7 Future Work and Directions

Our Neural CSSR approach opens several promising research directions that could significantly advance epsilon machine discovery and our understanding of neural sequence modeling.

7.1 Transfer Learning and Cross-Process Generalization

Multi-Process Training

Training transformers on multiple sequential processes simultaneously could lead to more robust and generalizable causal state representations:

- **Multi-task learning**: Train a single model on Golden Mean, Even Process, and Seven-State Human data to learn shared structural patterns
- **Internal representation analysis**: Investigate how exposure to diverse processes shapes the model's hidden representations and whether universal causal features emerge
- **Cross-process transfer**: Evaluate whether models trained on simpler processes (Golden Mean) can be fine-tuned more effectively on complex processes (Seven-State Human)

7.2 In-Context Learning for Epsilon Machine Discovery

ICL-Based State Discovery

Large language models demonstrate remarkable in-context learning capabilities that could revolutionize epsilon machine discovery:

- **Few-shot process identification**: Provide the model with a few examples of sequence-state pairs and evaluate its ability to infer causal states for new sequences
- **Zero-shot generalization**: Test whether models trained on one class of processes can discover states in completely novel process types without additional training

7.3 Representation Learning and Interpretability

Building on the inductive bias probing framework developed in our experiments, future work should explore:

- **Multitask representation shaping**: Design auxiliary tasks that encourage the formation of interpretable causal state representations in hidden layers

8 Multitask Latent Factorization Experiments

Objective

Shape nanoGPT's hidden representations so that ε-states and structural primitives become linearly accessible while preserving next-token performance. We use standard cross-entropy heads attached to selected layers (no custom losses).

8.1 Tasks and Setups

- **LM (baseline only)**: next-token prediction.
- **ε-state head**: per-token causal-state classification.
- **Distance head**: for the current ε-state, predict the shortest hop-count to every other state (6 classes including "unreachable").
- **Taps**: primarily `lm_head_input` (also tested early blocks).

8.2 Quantitative Summary (`lm_head_input`)

Configuration	Val LM loss	ε acc	Dist acc
Baseline (30k, frozen probes)	–	$\approx 89.3\%$	$\approx 32.9\%$
Baseline-35k (10k total LM)	≈ 0.540	$\approx 98.1\%$	$\approx 32.9\%$
Multitask (5k ε+dist)	≈ 0.540	$\approx 97.7\%$	$\approx 98.5\%$

Takeaway: LM-only training does not make structural distances linearly readable; multitask supervision dramatically improves factorization without harming LM loss.

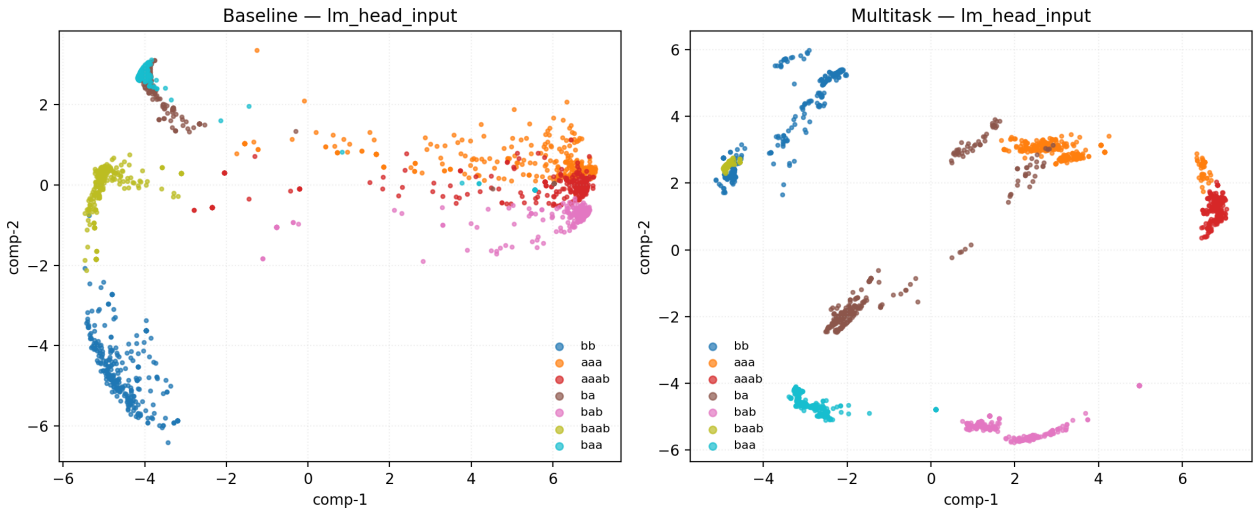

Figure 1: PCA at `lm_head_input`: Baseline vs. ε+Distance multitask. Multitask yields tight, margin-separated clusters with graph-consistent geometry.

8.3 Latent Projections

8.4 Current Explorations

- **Ablations**: (i) ε-only, (ii) distance-only, (iii) layer taps and bottlenecks.
- **Evaluation**: FER/NLI/fragmentation across layers; few-shot probe curves.
- **Data efficiency**: supervision budgets for heads vs. LM losses.

8.5 Configuration Pointers

- Baseline: `nanoGPT/out-seven-state-human-mtl100k-baseline/ckpt.pt`
- Multitask (ε+dist): `nanoGPT/mtl/config/train_multitask_seven_state_mtl100k_eps_dist.py`
- Distance-only: `nanoGPT/mtl/config/train_multitask_seven_state_mtl100k_dist_only.py`

9 OOD Robustness and Cheap Adaptation

9.1 Setup and Evaluator

Evaluator

We evaluate baseline vs. multitask finetunes (ε; ε+distance) under controlled OOD regimes using: `./scripts/eval_ood.sh` (wraps `nanoGPT/mtl/eval_ood.py`). Key flags:

- `--regime` in {`emission_mix`, `start_state_uniform`, `transition_noise`, `alphabet_swap`, `state_dependent_swap`, `temporal_swap`}
- `--remap-swap` or `--remap-matrix a,b,c,d`: post-hoc 2×2 probability remap (binary vocab)
- `--context-tokens K`: prefix-only scoring (few-shot support; score the tail)
- `--fit-remap {global,state}`: learn remap from the prefix (state = per-ε-state)

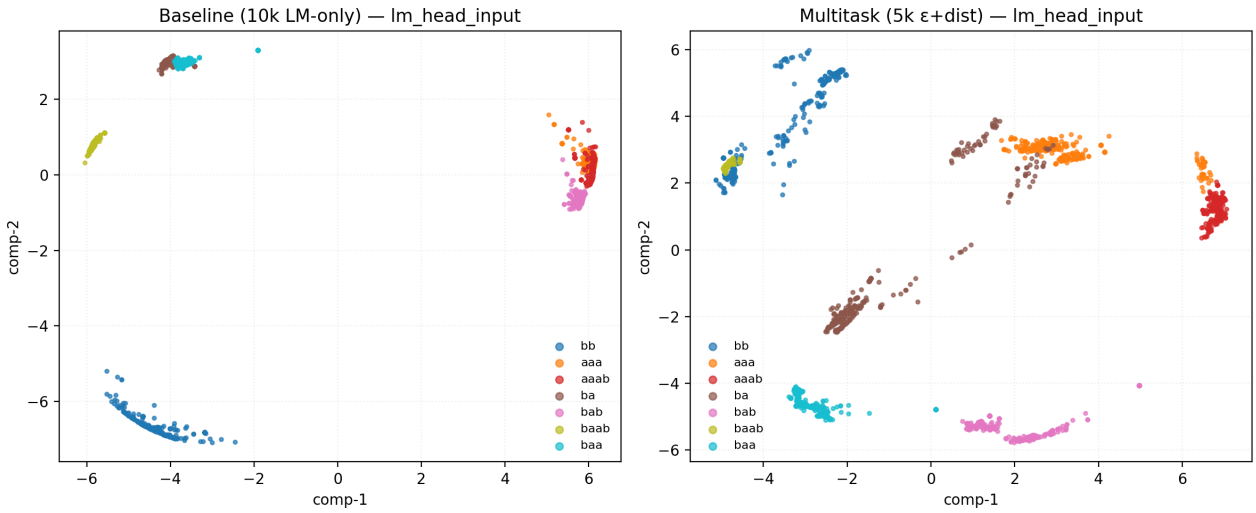

Figure 2: PCA comparison with explicit step labels: Baseline 10k LM-only vs. Multitask 5k ε+Distance. Extra LM steps sharpen ε separability but do not recover distance structure.

9.1.1 OOD Regimes

- **Emission mix** (α): soften emissions towards uniform (moderate covariate shift).
- **Start-state uniform**: randomize the initial causal state (tests recovery from different starts).
- **Transition noise** (p): random state jumps (dynamics perturbation).
- **Alphabet swap**: global permutation of 0/1 emissions (severe head miscalibration).
- **State-dependent swap** (fraction): swap emissions for a subset of states (head miscalibration conditioned on hidden state).
- **Temporal swap** (p): at each timestep, swap with probability p (time-varying mapping).

9.1.2 Protocol and Metrics

- **Token windows**: non-overlapping windows of size 64; total tokens per run: 65,536 unless noted.
- **Prefix-only scoring**: with `--context-tokens K`, only the suffix $t \in [K, \ldots, T)$ contributes to metrics; the prefix acts as observed support.
- **Head calibration**: `--remap-swap` applies a fixed 2×2 swap; `--fit-remap global` learns a single 2×2 matrix via least-squares on the prefix; `--fit-remap state` learns one 2×2 per ε-state using prefix labels.
- **Reported**: LM cross-entropy (bits/token), ε-accuracy, distance-accuracy (geometry), and scored token count.

9.2 Findings

Summary

- **Zero-shot LM**: On moderate shifts, LM ties; on severe emission permutation (alphabet swap), baseline has best raw LM. Multitask retains superior ε/geometry accuracy (structure remains aligned).

- **Cheap adaptation**: Tiny head remaps recover LM quickly. With `--remap-swap`, ε/ε+dist match/beat baseline on LM while keeping structural advantages.
- **Few-shot**: With `--context-tokens 16`, calibrated heads still favor ε+dist; context-fitted remaps (global/per-state) further reduce LM under harder regimes.
- **Hard regimes**: `state_dependent_swap` and `temporal_swap` break global remaps. Per-state fitted remaps (using ε-labels) restore LM to $\tilde{0}.61$ bits; temporal global fit drops LM to $\tilde{1}.00$ bits.

9.3 Representative Results

Alphabet Swap

Raw (zero-shot): `results/alphabet_swap_raw.csv:2-4`
Baseline 2.656, ε 2.735, ε+dist 2.817 bits/token.
Calibrated (`--remap-swap`): `results/alphabet_swap_calibrated.csv:2-4`
Baseline 0.539, ε 0.527, ε+dist 0.523 bits/token.
Few-shot + calibrated (`--context-tokens 16`): `results/alphabet_swap_context16_calibrated.csv:2-4`
Baseline 0.548, ε 0.535, ε+dist 0.531 bits/token.

State-Dependent Swap (harder)

Raw: `results/state_dependent_swap_0p5.csv:2-4`
Baseline 2.862, ε 2.975, ε+dist 3.067 bits/token; ε+dist keeps high distance acc ($\tilde{0}.678$).
Per-state fitted remap (`--context-tokens 16 --fit-remap state`): `results/state_dependent_swap_0p5_fitstate.csv:2-4`
Baseline 0.610, ε 0.641, ε+dist 0.618 bits/token.

Temporal Swap (time-varying)

Raw: `results/temporal_swap_0p5.csv:2-4` all $\tilde{1}.34$ bits/token.
Global fit from 16 support tokens: `results/temporal_swap_0p5_fitglobal.csv:2-4` all $\tilde{1}.00$ bits/token.

9.3.1 Compact Tables

Alphabet swap	Baseline	ε	ε+dist
Raw (bits/token)	2.656	2.735	2.817
Calibrated	0.539	0.527	**0.523**
Calibrated (ctx16)	0.548	0.535	**0.531**

State-dep swap (0.5)	Baseline	ε	ε+dist
Raw (bits/token)	**2.862**	2.975	3.067
Per-state fit (ctx16)	**0.610**	0.641	0.618

Temporal swap (0.5)	Baseline	ε	ε+dist
Raw (bits/token)	1.343	**1.334**	1.337
Global fit (ctx16)	**1.001**	1.001	1.001

9.4 Interpretation

Why Multitask Helps

- **Causal alignment**: Latents track ε-states/geometry; OOD head miscalibration is isolated and cheap to fix.
- **Linear separability**: Factorized latents enable tiny remaps/probes to work with few samples.
- **Failure isolation**: ε/geometry stable while LM degrades $\rightarrow$ diagnose and adapt the head only.
- **Compositional generalization**: Stable state geometry supports multi-step rollouts and state-aware adapters.

9.5 Methodological Notes and Limitations

- **Prefix fitting**: Current per-state fitting uses ground-truth ε on the prefix to establish an upper bound. In deployment, use ε-head predictions and confidence gating.
- **Binary vocab**: Remap examples are 2×2 (binary). For larger alphabets, use permutation matrices or low-rank adapters on logits.
- **What calibration changes**: Only LM loss changes under remaps; ε/distance metrics remain unchanged by design.
- **When remaps fail**: State/time-dependent mappings (hard regimes) require per-state or time-aware adapters; global remaps are insufficient.

9.6 Reproduction Snippets

Selected Commands

```
# Alphabet swap: raw vs calibrated vs few-shot calibrated
./scripts/eval_ood.sh --model baseline=nanoGPT/out-seven-state-human-mtl100k-baseline/ckpt.pt
    \
  --model eps=out-seven-state-human-mtl100k-mtl-eps/ckpt_final.pt \
  --model epsdist=out-seven-state-human-mtl100k-mtl-eps-dist/ckpt_final.pt \
  --regime alphabet_swap --severity 1.0 --tokens 65536 --output results/alphabet_swap_raw.csv
./scripts/eval_ood.sh ... --regime alphabet_swap --severity 1.0 --tokens 65536 \
  --remap-swap --output results/alphabet_swap_calibrated.csv
./scripts/eval_ood.sh ... --regime alphabet_swap --severity 1.0 --tokens 65536 \
  --context-tokens 16 --remap-swap --output results/alphabet_swap_context16_calibrated.csv

# Hard regime: per-state fitted remap from 16-token context
./scripts/eval_ood.sh ... --regime state_dependent_swap --severity 0.5 --tokens 65536 \
  --context-tokens 16 --fit-remap state --output results/state_dependent_swap_0p5_fitstate.
    csv

# Temporal swap: global fit from 16-token context
./scripts/eval_ood.sh ... --regime temporal_swap --severity 0.5 --tokens 65536 \
  --context-tokens 16 --fit-remap global --output results/temporal_swap_0p5_fitglobal.csv
```

10 Scaling to Real-World Sequential Data

The ultimate goal is extending Neural CSSR to natural sequential processes:

Real-World Applications

- **Time series analysis**: Apply to financial, biological, and physical time series data
- **Behavioral modeling**: Analyze sequential decision-making and learning processes
- **Computational neuroscience**: Model neural spike trains and brain state dynamics

These directions represent a natural evolution from our current binary process focus toward more general and impactful applications of neural-guided causal structure discovery.